

DATA 3464: Fundamentals of Data Processing

Wrangling text

Charlotte Curtis

March 5, 2026

Topic overview

- Text file formats and character encodings
- Extracting features from (un)structured text
- Regular expressions

Resources used:

- [Python documentation](#)
- [Joel on Software Blog Post](#)

What is a text file?

- Any file is just a sequence of **bytes**
- The file **extension** is somewhat meaningless
- Text files contain only human-readable **characters**
- **Binary files** are everything else, including:
 - Images
 - PDFs
 - Word documents*
 - Executables

ASCII

- American Standard Code for Information Interchange (1963)
- 7-bit encoding (128 characters)
- First 32 are **control characters** (e.g., newline, tab)
- Then punctuation, digits, uppercase letters, lowercase letters
- Most computers use 8-bit bytes, so there's a whole 128 characters "left over"


This is a very English-centric character set!

HOW STANDARDS PROLIFERATE:

(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)

SITUATION:
THERE ARE
14 COMPETING
STANDARDS.

14?! RIDICULOUS!
WE NEED TO DEVELOP
ONE UNIVERSAL STANDARD
THAT COVERS EVERYONE'S
USE CASES.



SOON:

SITUATION:
THERE ARE
15 COMPETING
STANDARDS.

Unicode to the rescue

- In the 1980s things were already getting out of hand
- The **Unicode consortium** published a standard in 1993 that assigned a **code point** to every character they could think of (297,334 as of Unicode 17.0)
- Currently, the most common encoding is **UTF-8**
 - "Unicode Transformation Format – 8-bit"
 - How can 297k characters be represented in only 8 bits?
- The first 128 characters are 1 byte each and align with ASCII
- After that, **2-4 bytes per character** are used

*There's also UTF-16 and UTF-32, but now we need to deal with **endianness***

Line endings

- Say we agree to use UTF-8 👍
- A text file is still just a big long string of **bytes**
- Humans like things to be orderly and readable with **line breaks**

System	Abbreviation	Escape sequence	Code point
Pre-OS X Mac	CR (carriage return)	<code>\r</code>	U+000D
Unix	LF (line feed)	<code>\n</code>	U+000A
Windows	CRLF	<code>\r\n</code>	U+000D U+000A

- Result: chaos, but we've mostly settled on LF (`\n`)

Why am I telling you all this?

- As data scientists, you will need to ingest data from various sources
- You **will** encounter character encoding and/or line ending issues
- You don't need to memorize all this, but recognizing that an issue exists will go a long ways towards fixing it

Example: misadventures with "smart" quotes

Where we left off on March 5

Portable document format

- The [PDF specification](#) was first published in 1993
- PDFs are like a "digital paper" that can contain text, images, vector graphics, and more (like JavaScript, forms, and weirdly, 3D models)
- We use them for all sorts of things because the **appearance** is consistent
- They are absolutely terrible for pretty much everything else

Some useful PDF packages:

- [pdfminer.six](#): extract text and metadata
- [pdfplumber](#): built on pdfminer.six, layout aware text extraction
- [pikepdf](#): low level PDF manipulation
- [pymupdf](#): more low level PDF manipulation
- and [many](#) ([many](#)) [more](#)!

DATA 3463 covers more about PDFs, including OCR. The main focus in this class is on fixing the issues and dealing with edge case

Hypertext markup language

- Much simpler than PDFs, but still not great for data exchange
- Markup languages are structured text files that define how content *should* be displayed, but it's not always consistent
- Most of the internet is UTF-8 encoded HTML
- We can try to extract text with something like BeautifulSoup
 - Again, more on web scraping in DATA 3463

Back to text

- Assuming we've got text from somewhere, we probably want to:
 - Organize it into a structured format (csv, database, json)
 - Identify specific features (postal codes, dates, names)
- Need to deal with encoding issues, garbled sentences, mixed up tables, and all the other bizarre ways things go wrong
- There is no magic flowchart for this! Lots of trying things, seeing what happens, dealing with a few edge cases at a time

*One really useful tool is **regular expressions** 🛠️*

Regular expressions

The card game?

Note: card groupings and colours are logical where possible, but sometimes just random

Basic characters

A literal character
e.g. A, x, 4, ?

\

Escape character

\d

Any digit

"Word" characters

\W

Any of
_, Aa-Zz, 0-9
(word characters)

\W

Any
non-word character

Whitespace

\s

Any
whitespace

\S

Any
non-whitespace

Quantifiers

*

0 or more

+

1 or more

?

0 or 1

Brackets and braces

[]

One of the
characters in []

(?:)

Exact sequence
in () after :

{#}

Match previous
exactly # times

Range and boundary

—

Range separator
e.g. A-Z

\b

Word boundary

\B

Not a
word boundary

Start (not) and end

^

Inside []: Not
Otherwise:
Start of line/string

\$

End of line/string

Where we left off on March 10

Capturing groups

- Seems weird that `(?:...)` is so verbose, why not just `(...)?`
- Even more common than grouping is **capturing**: `(...)`
- This lets us save part of the match for later

Examples:

- *extract the dollar amount from a \$\$ string*
- *extract the area code from a phone number*
- *extract the domain from an email address*

Regex tools

- [Regex101](#) for quickly testing, debugging, and learning
- [re module](#) for using regexes in Python
- [grep](#) for searching through files on the command line
- [sed](#) for doing regex-based find and replace on the command line
- Regex the card game for learning and fun?

Case study: Cloudflare outage

In July 2019, a poorly formed regex in Cloudflare's firewall rules caused CPU usage to spike and websites to come crashing down worldwide. The **cause**? This regex:

```
(?:(?:"|'|\)|\}|\\|\\d|(?:nan|infinity|true|false|null|undefined|symbol|math)|\\`|-|\\+|[+])*(?((?:\\$|-|~|!|{|}\\|\\\\|\\+)*.(?:.*=.*)))
```

- Regex engines perform **backtracking** to check for multiple match possibilities
- This particular string led to **catastrophic backtracking**
- Moral of the story: test carefully, both positive and negative

Coming up next

- Assignment 3: curate a dataset
- A primer on data cards
- Intro to signals and images