

DATA 3464: Fundamentals of Data Processing

Image Processing

Charlotte Curtis

March 24, 2026

Topic overview

- Image representation
- File formats and compression
- Preprocessing for image recognition tasks

Resources used:

- [Computer Vision: Algorithms and Applications \(2nd edition\)](#)
- [Pillow Documentation](#)
- [A blog post on image compression](#)
- [JPEG demo](#)

Images as 2D signals

- The light that enters a camera can be modelled as continuous signal:

$$f(x, y), \quad -\infty < x, y < \infty$$

- Digital images are sampled:

$$f[n, m], \quad n = n\Delta_x, m = m\Delta_y$$

where typically $\Delta_x = \Delta_y$

- The area $\Delta_x \times \Delta_y$ is called a **picture element**, or **pixel**





Image channels

- A photosensor responds to **light intensity** with an electrical signal
- Physical filters restrict the colour that reaches each sensor
- On most digital cameras, 1 sensor $\neq$ 1 pixel; it is instead **interpolated** to create a typical **3-channel image**
- ([Inexplicably](#), the Python imaging library Pillow uses the term "band")



Images as (3D) 2D arrays

- Size (`np.shape`) usually defined as `[rows, cols, (channels)]`
- Typical 8-bit image: integers from 0 to 255
 - For processing, common to convert to float and **normalize**
- RGB examples of individual pixels:
 - `[255, 0, 0]`
 - `[128, 128, 128]`
 - `[255, 255, 255]`

Portable Network Graphics

- **PNG** is an image format released in 1996 as a GIF replacement
- Typically 8 bits per channel, but can be 1, 2, 4, 8, or 16
- 3x RGB channels, plus an optional alpha (transparency) channel
- Space saved in two ways:
 - Indexed colour, where numbers map to a finite set of colours
 - Lossless compression, e.g. `[0 0 0 0 0 0 0]` → `7 0`

What kinds of images would benefit from this compression scheme?

Joint Photographic Experts Group

- **JPEG** is an image standard released in 1992 for photographic images
- **Lossy** compression based on human vision:
 - i. Channels converted to **YCbCr** colourspace
 - ii. The colour components (Cb/Cr) are downsampled 2-3x
 - iii. Blocks of 8x8 pixels are transformed to frequency domain
 - iv. Amplitudes are **quantized** with larger buckets at higher frequencies
 - v. Finally, lossless compression is applied
- The "quality" of a JPEG relates the degree of quantization

Where we left off on March 24

A Cautionary Tale of Image Classification

- As convolutional neural networks (CNNs) gained popularity as image recognition powerhouses, people started using them for all sorts of things
- One example: predicting whether someone is a criminal based on a photo
- **Thoroughly debunked** in a paper that mentions, among other issues:
 - Criminal samples are all 8-bit greyscale PNG photographs of printed images, taken with the same camera
 - Non-criminals are RGB JPEGs from 5 different sources, converted to greyscale to be "compatible with mugshots from the criminal category"

What's the problem here?